

Emotion Detection in Persian Text: A Classical Machine-Learning Baseline

Mahla Entezari

Shahid Beheshti University – Department of Computer Science

June 2024

Abstract

Emotion recognition from natural language is a foundational step toward empathetic human–computer interaction. This report presents a *classical* (non-neural) machine-learning pipeline that identifies five emotions (*happiness*, *sadness*, *anger*, *fear*, and *other*) in Persian social-media posts. The study covers data cleaning, exploratory data analysis (EDA), TF–IDF feature engineering, baseline logistic-regression training, boosting with XGBoost, thorough evaluation, and an error analysis that highlights class imbalance and linguistic subtleties. While deep transformers currently yield state-of-the-art performance in many languages, establishing a strong classical baseline is crucial for future comparative research and for low-resource environments where computation or annotation budgets are tight.

1 Introduction

Natural-language emotion detection enables a variety of downstream applications—ranging from content moderation to psychological well-being support—by automatically inferring users’ affective states. Although Persian (Farsi) is among the world’s top-15 most-spoken languages, publicly available emotion datasets and pretrained models remain scarce compared with English. Consequently, the pedagogical objective of the third assignment in the *Machine Learning Fundamentals* course is twofold:

- (i) design an end-to-end pipeline for emotion classification in Persian, and

- (ii) document results in a reproducible scientific report.

The rest of the manuscript is organized as follows: Section 2 describes the dataset and preprocessing; Section 3 summarizes exploratory findings; Section 4 details the modelling choices; Section 5 presents quantitative and qualitative results; Section 6 discusses limitations and potential improvements; finally, Section 7 concludes the work.

2 Dataset and Pre-processing

2.1 Raw Corpus

The provided corpus contains $\approx 4,200$ Persian sentences manually annotated with exactly one of five emotions. A separate unlabeled test set is reserved for blind evaluation on the course platform.

2.2 Normalization and Cleaning

Normalization is critical in Persian due to multiple Unicode variants of the same character. The following steps were scripted in `Python` (library versions in Appendix A):

- (1) Lower-case text; remove URLs, HTML entities, digits, and punctuation.
- (2) Tokenize with the `Hazm` tokenizer.
- (3) Remove 2-, 3-, and 4-grams of Persian stop-words compiled from the *Hamshahri* corpus.
- (4) Rejoin tokens into cleaned sentences for downstream vectorization.

2.3 Feature Engineering

Classical algorithms require fixed-length vectors. The cleaned texts were transformed with **TF-IDF** (*term-frequency* \times *inverse-document-frequency*). Label strings were mapped to integers via `sklearn.preprocessing.LabelEncoder`.

3 Exploratory Data Analysis

3.1 Class Distribution

Figure 1 reveals a pronounced imbalance: **HAPPY** and **OTHER** together account for nearly 60 % of examples, whereas **FEAR** barely surpasses 5 %. Such skew can inflate accuracy but harm minority-class recall, motivating stratified cross-validation and macro-averaged metrics.




Figure 1: Label frequencies in the training split. Color shading alternates only for readability.

3.2 Lexical Statistics

The most common tokens across all documents (occurring >100 times) are drawn in Figure 2. Many high-frequency words are topic neutral—e.g., generic verbs confirming the necessity of stop-word removal and bigram features to capture discriminative phrases.




Figure 2: Token frequency with logarithmic y-axis truncation for clarity.

4 Modelling Methodology

4.1 Train/Validation Protocol

- The corpus was stratified into an 80 % training and 20 % internal test split.
- All hyper-parameters were tuned via *stratified 5-fold* cross-validation on the training part.
- Final reported numbers on the internal test set proxy the hidden course evaluation.

4.2 Baseline Classifier: Logistic Regression

Logistic regression with ℓ_2 regularization is a strong linear baseline on sparse TF-IDF matrices. A grid search over $C \in \{0.01, 0.1, 1, 10\}$ favored $C = 1$.

4.3 Ensemble Extension: XGBoost

To gauge non-linear benefits, we fit an *Extreme Gradient Boosting* (XGBoost) model on the same TF-IDF vectors. Key parameters:

Parameter	Value
max_depth	: 6
learning_rate	: 0.1
n_estimators	: 100
subsample	: 0.8
colsample_bytree	: 0.8
objective	: multi:softprob

Training and validation error curves across 100 boosting rounds are plotted in Figure 3; early stopping after round 67 minimized validation loss.




Figure 3: Learning curves (mean \pm sd over CV folds) highlight diminishing returns and slight over-fit.

5 Results

5.1 Quantitative Evaluation

Table 1 contrasts logistic regression with XGBoost on the internal test set. All metrics are macro-averaged to reward balanced performance.

Table 1: Macro-averaged metrics on the held-out split.

Model	Precision	Recall	F1-score
Logistic Regression	0.20	0.21	0.11
XGBoost	0.42	0.38	0.37

Take-away. Boosting more than triples the macro-F1 compared with the linear baseline, evidencing non-linear separability among emotion classes in the TF-IDF space.

5.2 Class-Wise Breakdown

Table 2 drills into per-class performance for XGBoost only (best model). Notably, FEAR remains challenging due to scant data.

Table 2: Per-class scores for XGBoost.

Label	Precision	Recall	F1	Support
ANGRY	0.33	0.44	0.38	185
FEAR	0.21	0.17	0.19	66
HAPPY	0.66	0.74	0.70	306
OTHER	0.46	0.34	0.39	267
SAD	0.43	0.21	0.28	161

6 Discussion

6.1 Why Does Boosting Help?

Gradient-boosted trees exploit feature interactions ignored by linear methods.

6.2 Limitations

Imbalanced Data. Minority classes cap recall performance. Synthetic oversampling (e.g. SMOTE) or *class-weighted* loss in XGBoost warrants exploration.

Sparse Features. TF-IDF disregards word order; sentiment shifters such as negations may invert emotion yet share tokens.

6.3 Opportunities for Improvement

1. Fine-tune a multilingual transformer (mBERT, XLM-R) on the same labels.
2. Incorporate emoji embeddings or replace them with textual emotion synonyms before vectorization.
3. Apply *curriculum learning*: start with binary polarity, then expand to five-way emotion.

7 Conclusion

This assignment delivered a transparent yet non-trivial baseline for Persian emotion detection: data cleaning, EDA, TF-IDF vectorization, logistic and boosted classifiers, and in-depth interpretation. Although neural architectures are the logical next step, our classical pipeline already surfaces practical challenges—particularly class imbalance and lexical ambiguity—that any subsequent system must tackle. All code, hyper-parameters, and figures are open-sourced in the accompanying Jupyter notebook to facilitate reproducibility and extension.

A Software Environment

- Python 3.11.4
- scikit-learn 1.5
- XGBoost 2.0
- Hazm 0.9.0

B Assignment Requirements Checklist

- Data cleaning ✓

- Feature engineering ✓
- Tree-based classifier exploration ✓
- Evaluation with stratified CV ✓
- Inference pipeline (see notebook) ✓